

# Capstone - Muhammad Anwer

August 6, 2024

## 1 Data Loading and Preprocessing

```
[ ]: import matplotlib.pyplot as plt
from torchvision import datasets, transforms
from torch.utils.data import DataLoader, random_split
import numpy as np
import torch
import pandas as pd
```

```
[ ]: seed = 42

if (torch.cuda.is_available()):
    device = torch.device("cuda")
    torch.cuda.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)
    print("Running on GPU")
else:
    device = torch.device("cpu")
    torch.manual_seed(seed)
    print("Running on CPU")
```

Running on GPU

```
[ ]: transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])

dataset_dir = './food-101'
full_dataset = datasets.Food101(root=dataset_dir, split='train', download=True,
    ↪transform=transform)
classes = pd.DataFrame(full_dataset.classes, columns=['class'])
num_images_per_class = len(full_dataset) // len(full_dataset.classes)

print(f"Number of classes: {len(full_dataset.classes)}")
```

```
print(f"Number of images: {len(full_dataset)}")
print(f"Number of images per class: {num_images_per_class}")
print("Full dataset downloaded successfully")
```

Downloading <https://data.vision.ee.ethz.ch/cvl/food-101.tar.gz> to  
./food-101/food-101.tar.gz

100%| | 4996278331/4996278331 [03:46<00:00, 22055098.77it/s]

Extracting ./food-101/food-101.tar.gz to ./food-101

Number of classes: 101

Number of images: 75750

Number of images per class: 750

Full dataset downloaded successfully

```
[ ]: print("All classes:")
      classes
```

All classes:

```
[ ]:      class
      0      apple_pie
      1      baby_back_ribs
      2      baklava
      3      beef_carpaccio
      4      beef_tartare
      ..      ...
     96      tacos
     97      takoyaki
     98      tiramisu
     99      tuna_tartare
    100      waffles
```

[101 rows x 1 columns]

```
[ ]: num_filtered_classes = 20

      classes_to_keep = full_dataset.classes[:num_filtered_classes]
      classes_to_keep_df = pd.DataFrame(classes_to_keep, columns=['class'])

      print("Filtered classes:")
      classes_to_keep_df
```

Filtered classes:

```
[ ]:      class
      0      apple_pie
      1      baby_back_ribs
      2      baklava
```

```

3      beef_carpaccio
4      beef_tartare
5      beet_salad
6      beignets
7      bibimbap
8      bread_pudding
9      breakfast_burrito
10     bruschetta
11     caesar_salad
12     cannoli
13     caprese_salad
14     carrot_cake
15     ceviche
16     cheese_plate
17     cheesecake
18     chicken_curry
19     chicken_quesadilla

```

```

[ ]: # Filter the dataset to keep only the first 20 classes
try:
    filtered_dataset = torch.load('filtered_dataset.pth')
    print("Filtered dataset loaded successfully")
except:
    indices_to_keep = [i for i, (_, class_idx) in enumerate(full_dataset) if
        ↪full_dataset.classes[class_idx] in classes_to_keep]
    filtered_dataset = torch.utils.data.Subset(full_dataset, indices_to_keep)

# Split the filtered data into training and validation sets
total_train_and_val_size = num_filtered_classes * num_images_per_class
train_size = int(total_train_and_val_size * 0.8)
val_size = int(total_train_and_val_size * 0.2)
train_dataset, val_dataset = random_split(filtered_dataset, [train_size,
    ↪val_size])

print(f"Number of training images: {len(train_dataset)}")
print(f"Number of validation images: {len(val_dataset)}")

```

```

Filtered dataset loaded successfully
Number of training images: 12000
Number of validation images: 3000

```

```

[ ]: # Filtering takes a long time, so save the filtered dataset
torch.save(filtered_dataset, 'filtered_dataset.pth')
print("Filtered dataset saved successfully")

```

```

Filtered dataset saved successfully

```

```
[ ]: def denormalize(tensor, mean, std):
    tensor = tensor.clone() # Avoid modifying the original tensor
    for t, m, s in zip(tensor, mean, std):
        t.mul_(s).add_(m)
    return tensor

# Visualize some samples
def visualize_samples(loader, num_samples=5):
    dataiter = iter(loader)
    images, labels = next(dataiter)
    images = images[:num_samples]
    labels = labels[:num_samples]

    plt.figure(figsize=(15, 5))
    for i in range(num_samples):
        image = denormalize(images[i], [0.485, 0.456, 0.406], [0.229, 0.224, 0.
↪225])
        image = image.permute(1, 2, 0).numpy() # Convert CHW to HWC
        plt.subplot(1, num_samples, i + 1)
        plt.imshow(np.clip(image, 0, 1))
        plt.title(full_dataset.classes[labels[i]])
        plt.axis('off')
    plt.show()
```

```
[ ]: def create_dataloaders(batch_size):
    trainloader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
    valloader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
    print("DataLoaders created successfully")
    return trainloader, valloader
```

```
[ ]: trainloader, valloader = create_dataloaders(64)

visualize_samples(trainloader, num_samples=5)
```

DataLoaders created successfully



```
[ ]: # Load the test dataset
test_dataset = datasets.Food101(root=dataset_dir, split='test', download=True,
    ↪transform=transform)
num_images_per_class_test = len(test_dataset) // len(test_dataset.classes)

print(f"Number of classes: {len(test_dataset.classes)}")
print(f"Number of images: {len(test_dataset)}")
print(f"Number of images per class: {num_images_per_class_test}")
print("Full dataset downloaded successfully")
```

Number of classes: 101  
 Number of images: 25250  
 Number of images per class: 250  
 Full dataset downloaded successfully

```
[ ]: try:
    filtered_test_dataset = torch.load('filtered_test_dataset.pth')
    print("Filtered test dataset loaded successfully")
except:
    test_indices_to_keep = [i for i, (_, class_idx) in enumerate(test_dataset) if
    ↪test_dataset.classes[class_idx] in classes_to_keep]
    filtered_test_dataset = torch.utils.data.Subset(test_dataset,
    ↪test_indices_to_keep)

test_size = 2000
filtered_test_dataset, _ = random_split(filtered_test_dataset, [test_size,
    ↪len(filtered_test_dataset) - test_size])
print(f"Number of test images: {len(filtered_test_dataset)}")
```

Filtered test dataset loaded successfully  
 Number of test images: 2000

```
[ ]: torch.save(filtered_test_dataset, 'filtered_test_dataset.pth')
print("Filtered test dataset saved successfully")
```

Filtered test dataset saved successfully

```
[ ]: testloader = DataLoader(filtered_test_dataset, batch_size=64, shuffle=False)
print("Test DataLoader created successfully")
```

Test DataLoader created successfully

## 2 Model

```
[ ]: import torch.nn as nn
import torch.optim as optim
from torchvision.models import resnet18, ResNet18_Weights
from tqdm import tqdm

def initialize_model(layers_to_unfreeze):
    # Load pre-trained model & freeze the layers
    model = resnet18(weights=ResNet18_Weights.DEFAULT)

    if layers_to_unfreeze == 0:
        print('Freeze all layers')
        for param in model.parameters():
            param.requires_grad = False
    elif layers_to_unfreeze < 18:
        print('Freeze all but the last n layers')
        for param in list(model.parameters())[:-layers_to_unfreeze]:
            param.requires_grad = False
    elif layers_to_unfreeze >= 18:
        print('Unfreeze all the layers')
        for param in model.parameters():
            param.requires_grad = True

    # Replace the last fully connected layer with a 20-classes layer (since we
    ↪ have chosen to filter out 20 classes)
    features = model.fc.in_features
    model.fc = nn.Sequential(
        nn.Dropout(0.5), # helped improve the loss
        nn.Linear(features, 20)
    )

    # Move the model to GPU if available
    if torch.cuda.is_available():
        device = torch.device("cuda")
        model = model.to(device)

    return model
```

```
[ ]: import torch.optim as optim
from tqdm import tqdm

def train_model(trainloader, valloader, lr, epochs=10, momentum=0,
    ↪ best_val_loss=1000.00, layers_to_unfreeze=0):
    model = initialize_model(layers_to_unfreeze)
    criterion = nn.CrossEntropyLoss()
```

```

# by default use Adam
if momentum == 0:
    optimizer = optim.Adam(model.fc.parameters(), lr=lr, weight_decay=1e-4)
else:
    optimizer = optim.SGD(model.parameters(), lr=lr, momentum=momentum,
↪weight_decay=5e-4)

train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []
iteration_losses = []
iteration_accuracies = []

for epoch in tqdm(range(epochs)):
    model.train()
    epoch_running_loss = 0.0
    epoch_correct = 0
    epoch_total = 0

    interval_running_loss = 0.0
    interval_correct = 0
    interval_total = 0

    for i, data in enumerate(trainloader, 0):
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()

        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        epoch_running_loss += loss.item()
        interval_running_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        epoch_total += labels.size(0)
        epoch_correct += (predicted == labels).sum().item()

        interval_total += labels.size(0)
        interval_correct += (predicted == labels).sum().item()

    if i % 100 == 99:
        iteration_loss = interval_running_loss / 100
        iteration_accuracy = 100 * interval_correct / interval_total

```

```

        print(f'[Epoch {epoch+1}, Iteration {i+1}] Training loss:␣
↪{iteration_loss:.4f}, Training accuracy: {iteration_accuracy:.2f}%')
        iteration_losses.append(iteration_loss)
        iteration_accuracies.append(iteration_accuracy)
        interval_running_loss = 0.0
        interval_correct = 0
        interval_total = 0

epoch_loss = epoch_running_loss / len(trainloader)
epoch_accuracy = 100 * epoch_correct / epoch_total
train_losses.append(epoch_loss)
train_accuracies.append(epoch_accuracy)

# Validation at the end of the epoch
model.eval()
val_loss = 0.0
val_correct = 0
val_total = 0

with torch.no_grad():
    for inputs, labels in valloader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        val_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        val_total += labels.size(0)
        val_correct += (predicted == labels).sum().item()

epoch_val_loss = val_loss / len(valloader)
epoch_val_accuracy = 100 * val_correct / val_total
val_losses.append(epoch_val_loss)
val_accuracies.append(epoch_val_accuracy)
print(f'\n>> Epoch {epoch+1}, Validation loss: {epoch_val_loss:.4f},␣
↪Validation accuracy: {epoch_val_accuracy:.2f}%')

if epoch_val_loss < best_val_loss:
    print(f'>>>> New best model found with validation loss:␣
↪{epoch_val_loss:.4f}')
    print(f'>>>> Learning rate: {lr}, Epochs: {epochs}')
    if momentum == 0:
        print(f'>>>> Optimizer: Adam')
    else:
        print(f'>>>> Optimizer: SGD with momentum {momentum}')
    print(f'>>>> Layers unfrozen: {layers_to_unfreeze}')
    best_val_loss = epoch_val_loss
    torch.save(model.state_dict(), 'best_model.pth')

```



```

    return train_losses, val_losses, train_accuracies, val_accuracies,
    ↪ best_val_loss, iteration_losses, iteration_accuracies

```

## 3 Training

### 3.1 First Attempt

```

[ ]: # from tqdm import tqdm

# train_losses = []
# train_accuracy = []
# val_losses = []
# val_accuracy = []

# optimizer = optim.Adam(model.fc.parameters(), lr=0.001)
# criterion = nn.CrossEntropyLoss()

# model = resnet18(weights=ResNet18_Weights.DEFAULT)

# for param in model.parameters(): # Freeze all but the last 5 layers
#     param.requires_grad = False

# # Replace the last fully connected layer with a 20-classes layer
# features = model.fc.in_features
# model.fc = nn.Linear(features, 20)

# # Move the model to GPU if available
# if torch.cuda.is_available():
#     device = torch.device("cuda")
#     model = model.to(device)

# # Train the model - write the training loop
# '''
# # in the loop you need to keep track of training and validation accuracy.
# # Print the training & validation loss & accuracy after every epoch.

# # Remember when you evaluate the model on val or test set, set the model to eval
# # mode, and turn off gradients requirement. Once you're done, set the model back
# ↪ to
# train mode
# '''
# for epoch in tqdm(range(10)): # loop over the dataset multiple times

#     running_loss = 0.0

```

```

#     correct = 0
#     total = 0
#     for i, data in enumerate(trainloader, 0):
#         # get the inputs; data is a list of [inputs, labels]
#         inputs, labels = data
#         inputs, labels = inputs.to(device), labels.to(device)

#         # zero the parameter gradients
#         optimizer.zero_grad()

#         # forward + backward + optimize
#         outputs = model(inputs)
#         loss = criterion(outputs, labels)
#         loss.backward()
#         optimizer.step()

#         # print statistics
#         running_loss += loss.item()

#         # calculate training accuracy
#         _, predicted = torch.max(outputs, 1)
#         total += labels.size(0)
#         correct += (predicted == labels).sum().item()

#         # print & accumulate training statistics (loss & accuracy) every 100
#         ↪ iterations
#         if i % 100 == 99:
#             print(f'[Epoch {epoch+1}, Iteration {i+1}] loss: {running_loss /
#             ↪ 100:.4f}, accuracy: {100 * correct / total:.2f}%')
#             train_losses.append(running_loss / 100)
#             train_accuracy.append(100 * correct / total)
#             running_loss = 0.0
#             correct = 0
#             total = 0

#     # train_loss = running_loss / len(trainloader)
#     # train_acc = 100 * correct / total
#     # train_losses.append(train_loss)
#     # train_accuracy.append(train_acc)

#     # Validation at the end of the epoch
#     # set mode to eval mode & tell torch no gradients are required.
#     model.eval()

#     # define any valuable you need to track here <<<<<
#     val_loss = 0.0
#     val_correct = 0

```

```

#     val_total = 0

#     with torch.no_grad():

#         # loop through the data and find accuracy and loss for all val data
#         for data in valloader:
#             inputs, labels = data
#             inputs, labels = inputs.to(device), labels.to(device)
#             outputs = model(inputs)
#             loss = criterion(outputs, labels)
#             val_loss += loss.item()
#             _, predicted = torch.max(outputs, 1)
#             val_total += labels.size(0)
#             val_correct += (predicted == labels).sum().item()

#     epoch_val_loss = val_loss / len(valloader)
#     epoch_val_accuracy = 100 * val_correct / val_total
#     val_losses.append(epoch_val_loss)
#     val_accuracy.append(epoch_val_accuracy)

#     # set the model back to train mode
#     model.train()

# print('Finished Training')
# #

```

```

0%|          | 0/10 [00:00<?, ?it/s]
[Epoch 1, Iteration 100] loss: 2.4029, accuracy: 30.86%
10%|         | 1/10 [01:22<12:21, 82.37s/it]
[Epoch 2, Iteration 100] loss: 1.4533, accuracy: 60.12%
20%|         | 2/10 [02:43<10:54, 81.75s/it]
[Epoch 3, Iteration 100] loss: 1.2566, accuracy: 63.56%
30%|         | 3/10 [04:05<09:31, 81.60s/it]
[Epoch 4, Iteration 100] loss: 1.1635, accuracy: 65.98%
40%|         | 4/10 [05:28<08:12, 82.15s/it]
[Epoch 5, Iteration 100] loss: 1.0997, accuracy: 67.45%
50%|         | 5/10 [06:49<06:49, 81.90s/it]
[Epoch 6, Iteration 100] loss: 1.0577, accuracy: 68.27%
60%|         | 6/10 [08:10<05:26, 81.69s/it]

```

```
[Epoch 7, Iteration 100] loss: 1.0127, accuracy: 69.91%
70%|          | 7/10 [09:31<04:04, 81.41s/it]
[Epoch 8, Iteration 100] loss: 1.0004, accuracy: 70.00%
80%|          | 8/10 [10:53<02:43, 81.52s/it]
[Epoch 9, Iteration 100] loss: 0.9888, accuracy: 69.97%
90%|          | 9/10 [12:14<01:21, 81.51s/it]
[Epoch 10, Iteration 100] loss: 0.9748, accuracy: 70.91%
100%|         | 10/10 [13:35<00:00, 81.58s/it]
Finished Training
```

## 4 Hyperparameter Tuning

We will be tuning the following:

- learning rates
- batch sizes
- number of epochs
- SGD optimizer or Adam
- number of hidden layers to unfreeze in the model

I went back-and-forth a number of times. The version below is the final version of the function.

```
[ ]: # set a high best_val_loss value so that by default the model can take its best
      ↪ value and return it
def tune_model(learning_rates, batch_sizes, epochs_list, momentum_list,
      ↪ layers_to_unfreeze, best_val_loss=1000.00):

    # check if all of the lists that are passed as params are of the same length
    assert len(learning_rates) == len(batch_sizes) == len(epochs_list) ==
    ↪ len(momentum_list) == len(layers_to_unfreeze), "All lists must have the same
    ↪ length"

    experiment_results = []

    # Run experiments
    for index, lr in enumerate(learning_rates):
        batch_size = batch_sizes[index]
        epochs = epochs_list[index]
        momentum = momentum_list[index]
        layers_to_unfreeze_value = layers_to_unfreeze[index]

        print(f'Running experiment with lr={lr}, batch_size={batch_size},
        ↪ epochs={epochs}')
```

```

if momentum > 0:
    print(f'Using SGD with momentum: {momentum}')
else:
    print(f'Using Adam')
print(f'Unfreezing {layers_to_unfreeze_value} layers')

trainloader, valloader = create_dataloaders(batch_size)
train_losses, val_losses, train_accuracies, val_accuracies,
↪best_val_loss, iteration_losses, iteration_accuracies =
↪train_model(trainloader, valloader, lr, epochs, momentum, best_val_loss,
↪layers_to_unfreeze_value)
experiment_results.append({
    'lr': lr,
    'batch_size': batch_size,
    'epochs': epochs,
    'optimizer': 'SGD' if momentum > 0 else 'Adam',
    'momentum': momentum,
    'layers_to_unfreeze': layers_to_unfreeze_value,
    'train_losses': train_losses,
    'val_losses': val_losses,
    'train_accuracies': train_accuracies,
    'val_accuracies': val_accuracies,
    'iteration_losses': iteration_losses,
    'iteration_accuracies': iteration_accuracies,
    'best_val_loss': best_val_loss,
})
print('\n')

return experiment_results

```

#### 4.1 First Pass at Hyperparameter Tuning

```

[ ]: # verify that the hyperparamter tuning works correctly

learning_rates = [0.01]
batch_sizes = [32]
epochs_list = [3]
momentum_list = [0.5]
layers_to_unfreeze = [0] # all hidden layers are frozen

experiment_results = tune_model(learning_rates, batch_sizes, epochs_list,
↪momentum_list, layers_to_unfreeze)

```

Running experiment with lr=0.01, batch\_size=32, epochs=3  
Using SGD with momentum: 0.5  
Unfreezing 0 layers  
DataLoaders created successfully

Freeze all layers

```
0%|          | 0/3 [00:00<?, ?it/s]

[Epoch 1, Iteration 100] Training loss: 2.7936, Training accuracy: 17.16%
[Epoch 1, Iteration 200] Training loss: 2.2895, Training accuracy: 31.19%
[Epoch 1, Iteration 300] Training loss: 2.0212, Training accuracy: 39.84%

33%|          | 1/3 [01:12<02:24, 72.50s/it]

>> Epoch 1, Validation loss: 1.5502, Validation accuracy: 56.73%
>>>> New best model found with validation loss: 1.5502
>>>> Learning rate: 0.01, Epochs: 3
>>>> Optimizer: SGD with momentum 0.5
[Epoch 2, Iteration 100] Training loss: 1.8114, Training accuracy: 45.84%
[Epoch 2, Iteration 200] Training loss: 1.8099, Training accuracy: 45.22%
[Epoch 2, Iteration 300] Training loss: 1.8146, Training accuracy: 46.12%

67%|          | 2/3 [02:25<01:12, 72.70s/it]

>> Epoch 2, Validation loss: 1.4164, Validation accuracy: 58.73%
>>>> New best model found with validation loss: 1.4164
>>>> Learning rate: 0.01, Epochs: 3
>>>> Optimizer: SGD with momentum 0.5
[Epoch 3, Iteration 100] Training loss: 1.7439, Training accuracy: 47.16%
[Epoch 3, Iteration 200] Training loss: 1.7541, Training accuracy: 47.72%
[Epoch 3, Iteration 300] Training loss: 1.7529, Training accuracy: 48.38%

100%|         | 3/3 [03:38<00:00, 72.80s/it]

>> Epoch 3, Validation loss: 1.3736, Validation accuracy: 58.73%
>>>> New best model found with validation loss: 1.3736
>>>> Learning rate: 0.01, Epochs: 3
>>>> Optimizer: SGD with momentum 0.5
```

```
[ ]: experiment_results
```

```
[ ]: [{'lr': 0.01,
      'batch_size': 32,
      'epochs': 3,
      'optimizer': 'SGD',
      'momentum': 0.5,
      'layers_to_unfreeze': 0,
      'train_losses': [2.2758195323944093, 1.7994298833211264, 1.7531641569137573],
      'val_losses': [1.5501532592671983, 1.4163949172547523, 1.3735889044213803],
```

```

'train_accuracies': [32.1, 46.00833333333333, 47.583333333333336],
'val_accuracies': [56.733333333333334,
58.733333333333334,
58.733333333333334],
'iteration_losses': [2.7935571420192717,
2.289516499042511,
2.0211995029449463,
1.811358060836792,
1.8099000465869903,
1.8145918536186219,
1.7439234471321106,
1.754069937467575,
1.7528611433506012],
'iteration_accuracies': [17.15625,
31.1875,
39.84375,
45.84375,
45.21875,
46.125,
47.15625,
47.71875,
48.375],
'best_val_loss': 1.3735889044213803}]

```

## 4.2 More Comprehensive Tuning

```

[ ]: learning_rates = [0.01, 0.008, 0.001]
batch_sizes = [32, 64, 64]
epochs_list = [3, 3, 3]
momentum_list = [0.5, 0.4, 0] # 0 means use Adam Optimizer
layers_to_unfreeze = [18, 10, 10]

best_val = experiment_results[0]['best_val_loss']

experiment_results_2 = tune_model(learning_rates, batch_sizes, epochs_list,
↪momentum_list, layers_to_unfreeze, best_val)

```

Running experiment with lr=0.01, batch\_size=32, epochs=3

Using SGD with momentum: 0.5

Unfreezing 18 layers

DataLoaders created successfully

Unfreeze all the layers

```
0%|          | 0/3 [00:00<?, ?it/s]
```

[Epoch 1, Iteration 100] Training loss: 2.3120, Training accuracy: 30.66%

[Epoch 1, Iteration 200] Training loss: 1.4147, Training accuracy: 57.06%

[Epoch 1, Iteration 300] Training loss: 1.2280, Training accuracy: 62.75%

33%| | 1/3 [01:17<02:35, 77.88s/it]

```
>> Epoch 1, Validation loss: 1.0563, Validation accuracy: 67.07%
>>>> New best model found with validation loss: 1.0563
>>>> Learning rate: 0.01, Epochs: 3
>>>> Optimizer: SGD with momentum 0.5
[Epoch 2, Iteration 100] Training loss: 0.8324, Training accuracy: 74.62%
[Epoch 2, Iteration 200] Training loss: 0.8061, Training accuracy: 75.00%
[Epoch 2, Iteration 300] Training loss: 0.7749, Training accuracy: 75.53%
```

67%| | 2/3 [02:36<01:18, 78.07s/it]

```
>> Epoch 2, Validation loss: 0.9513, Validation accuracy: 71.07%
>>>> New best model found with validation loss: 0.9513
>>>> Learning rate: 0.01, Epochs: 3
>>>> Optimizer: SGD with momentum 0.5
[Epoch 3, Iteration 100] Training loss: 0.5025, Training accuracy: 84.50%
[Epoch 3, Iteration 200] Training loss: 0.5105, Training accuracy: 84.41%
[Epoch 3, Iteration 300] Training loss: 0.5279, Training accuracy: 83.88%
```

100%| | 3/3 [03:53<00:00, 77.98s/it]

```
>> Epoch 3, Validation loss: 0.7693, Validation accuracy: 76.87%
>>>> New best model found with validation loss: 0.7693
>>>> Learning rate: 0.01, Epochs: 3
>>>> Optimizer: SGD with momentum 0.5
```

Running experiment with lr=0.008, batch\_size=64, epochs=3  
Using SGD with momentum: 0.4  
Unfreezing 10 layers  
DataLoaders created successfully

Freeze all but the last n layers

0%| | 0/3 [00:00<?, ?it/s]

```
[Epoch 1, Iteration 100] Training loss: 2.6675, Training accuracy: 19.72%
```

33%| | 1/3 [01:15<02:30, 75.08s/it]

```
>> Epoch 1, Validation loss: 1.5369, Validation accuracy: 57.20%
[Epoch 2, Iteration 100] Training loss: 1.5752, Training accuracy: 53.03%
```

67%| | 2/3 [02:29<01:14, 74.98s/it]

```
>> Epoch 2, Validation loss: 1.2071, Validation accuracy: 64.33%
```



```
[Epoch 3, Iteration 100] Training loss: 1.2785, Training accuracy: 61.55%
100%|          | 3/3 [03:45<00:00, 75.01s/it]
```

```
>> Epoch 3, Validation loss: 1.1004, Validation accuracy: 66.50%
```

```
Running experiment with lr=0.001, batch_size=64, epochs=3
Using Adam
Unfreezing 10 layers
DataLoaders created successfully
Freeze all but the last n layers
```

```
0%|          | 0/3 [00:00<?, ?it/s]
[Epoch 1, Iteration 100] Training loss: 2.7307, Training accuracy: 18.59%
33%|          | 1/3 [01:14<02:28, 74.40s/it]
```

```
>> Epoch 1, Validation loss: 1.7566, Validation accuracy: 54.33%
[Epoch 2, Iteration 100] Training loss: 1.8801, Training accuracy: 44.28%
67%|          | 2/3 [02:28<01:14, 74.17s/it]
```

```
>> Epoch 2, Validation loss: 1.5055, Validation accuracy: 58.50%
[Epoch 3, Iteration 100] Training loss: 1.7252, Training accuracy: 48.41%
100%|          | 3/3 [03:43<00:00, 74.43s/it]
```

```
>> Epoch 3, Validation loss: 1.4056, Validation accuracy: 59.83%
```

```
[ ]: experiment_results_2
```

```
[ ]: [{'lr': 0.01,
      'batch_size': 32,
      'epochs': 3,
      'optimizer': 'SGD',
      'momentum': 0.5,
      'layers_to_unfreeze': 18,
      'train_losses': [1.5416994927724204, 0.8032998402913412, 0.5167996644377708],
      'val_losses': [1.0562903357947127, 0.9513036427979774, 0.7692997512665201],
      'train_accuracies': [53.675, 75.2, 83.99166666666666],
      'val_accuracies': [67.06666666666666, 71.06666666666666, 76.86666666666666],
      'iteration_losses': [2.311954411268234,
```

```

1.4146864902973175,
1.2280431115627288,
0.8323881134390831,
0.8060550811886787,
0.7748781517148018,
0.5024746052920819,
0.5105489372462034,
0.5278520260751247],
'iteration_accuracies': [30.65625,
57.0625,
62.75,
74.625,
75.0,
75.53125,
84.5,
84.40625,
83.875],
'best_val_loss': 0.7692997512665201},
{'lr': 0.008,
'batch_size': 64,
'epochs': 3,
'optimizer': 'SGD',
'momentum': 0.4,
'layers_to_unfreeze': 10,
'train_losses': [2.3239665240683456, 1.4926749277622142, 1.2464442177021757],
'val_losses': [1.536876120465867, 1.2071344801720152, 1.1003550912471527],
'train_accuracies': [30.308333333333334,
55.083333333333336,
62.34166666666667],
'val_accuracies': [57.2, 64.33333333333333, 66.5],
'iteration_losses': [2.667481096982956,
1.5751827621459962,
1.2785279923677444],
'iteration_accuracies': [19.71875, 53.03125, 61.546875],
'best_val_loss': 0.7692997512665201},
{'lr': 0.001,
'batch_size': 64,
'epochs': 3,
'optimizer': 'Adam',
'momentum': 0,
'layers_to_unfreeze': 10,
'train_losses': [2.4461656206465783, 1.8474177179184366, 1.71366190149429],
'val_losses': [1.7566089985218454, 1.5055467980973265, 1.405574823947663],
'train_accuracies': [27.141666666666666,
45.06666666666667,
48.608333333333334],
'val_accuracies': [54.333333333333336, 58.5, 59.833333333333336],

```

```
'iteration_losses': [2.730701208114624,
1.880145592689514,
1.7252239167690278],
'iteration_accuracies': [18.59375, 44.28125, 48.40625],
'best_val_loss': 0.7692997512665201}]
```

### 4.3 Visualize the Results

```
[ ]: def plot_combined_experiment_results(combined_experiment_results):
    epochs = list(range(1, combined_experiment_results[0]['epochs'] + 1))

    plt.figure(figsize=(10, 6))
    for idx, result in enumerate(combined_experiment_results):
        label = f"{result['optimizer']}, Layers:␣
↪{result['layers_to_unfreeze']}, LR: {result['lr']}"
        plt.plot(epochs, result['train_losses'], label=f'{label} - Train')
        plt.plot(epochs, result['val_losses'], linestyle='--', label=f'{label}␣
↪- Val')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.title('Training and Validation Loss')
    plt.legend()
    plt.show()

    plt.figure(figsize=(10, 6))
    for idx, result in enumerate(combined_experiment_results):
        label = f"{result['optimizer']}, Layers:␣
↪{result['layers_to_unfreeze']}, LR: {result['lr']}"
        plt.plot(epochs, result['train_accuracies'], label=f'{label} - Train')
        plt.plot(epochs, result['val_accuracies'], linestyle='--',␣
↪label=f'{label} - Val')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy (%)')
    plt.title('Training and Validation Accuracy')
    plt.legend()
    plt.show()

    plt.figure(figsize=(10, 6))
    for idx, result in enumerate(combined_experiment_results):
        label = f"Opt: {result['optimizer']}, Layers:␣
↪{result['layers_to_unfreeze']}, LR: {result['lr']}"
        plt.plot(result['iteration_losses'], label=f'{label} - Iteration Loss')
    plt.xlabel('Iterations (per 100)')
    plt.ylabel('Loss')
    plt.title('Iteration Loss')
    plt.legend()
```

```

plt.show()


plt.figure(figsize=(10, 6))
for idx, result in enumerate(combined_experiment_results):
    label = f"Opt: {result['optimizer']}, Layers: {
↳{result['layers_to_unfreeze']}, LR: {result['lr']}"
    plt.plot(result['iteration accuracies'], label=f'{label} - Iteration_
↳Accuracy')
plt.xlabel('Iterations (per 100)')
plt.ylabel('Accuracy (%)')
plt.title('Iteration Accuracy')
plt.legend()
plt.show()



```

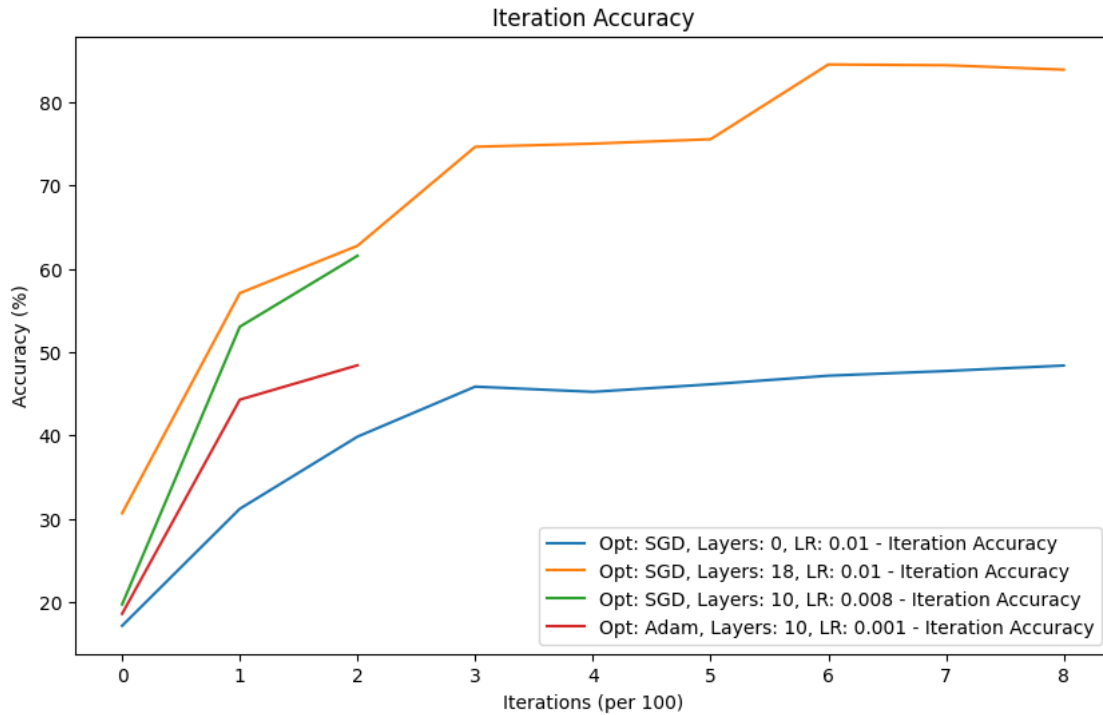
```

[ ]: combined_experiment_results = experiment_results + experiment_results_2
plot_combined_experiment_results(combined_experiment_results)

```







#### 4.4 Fine Tuning

```
[ ]: learning_rates = [0.01, 0.008, 0.001]
batch_sizes = [32, 64, 64]
epochs_list = [6, 10, 10] # increase epochs
momentum_list = [0.5, 0.4, 0]
layers_to_unfreeze = [18, 10, 10]

best_val = experiment_results_2[0]['best_val_loss']

experiment_results_3 = tune_model(learning_rates, batch_sizes, epochs_list,
↪momentum_list, layers_to_unfreeze, best_val)
```

Running experiment with lr=0.01, batch\_size=32, epochs=6

Using SGD with momentum: 0.5

Unfreezing 18 layers

DataLoaders created successfully

Unfreeze all the layers

0% | 0/6 [00:00<?, ?it/s]

[Epoch 1, Iteration 100] Training loss: 2.3233, Training accuracy: 31.22%

[Epoch 1, Iteration 200] Training loss: 1.4126, Training accuracy: 57.22%

[Epoch 1, Iteration 300] Training loss: 1.2352, Training accuracy: 63.16%

```

17%|          | 1/6 [01:19<06:36, 79.28s/it]

>> Epoch 1, Validation loss: 0.9603, Validation accuracy: 70.97%
[Epoch 2, Iteration 100] Training loss: 0.8188, Training accuracy: 76.00%
[Epoch 2, Iteration 200] Training loss: 0.8395, Training accuracy: 74.56%
[Epoch 2, Iteration 300] Training loss: 0.7884, Training accuracy: 76.44%
33%|          | 2/6 [02:38<05:16, 79.19s/it]

>> Epoch 2, Validation loss: 0.7896, Validation accuracy: 76.03%
[Epoch 3, Iteration 100] Training loss: 0.5118, Training accuracy: 83.97%
[Epoch 3, Iteration 200] Training loss: 0.5348, Training accuracy: 83.19%
[Epoch 3, Iteration 300] Training loss: 0.5533, Training accuracy: 82.66%
50%|          | 3/6 [03:57<03:57, 79.01s/it]

>> Epoch 3, Validation loss: 0.8599, Validation accuracy: 75.10%
[Epoch 4, Iteration 100] Training loss: 0.3089, Training accuracy: 90.38%
[Epoch 4, Iteration 200] Training loss: 0.3057, Training accuracy: 90.88%
[Epoch 4, Iteration 300] Training loss: 0.3250, Training accuracy: 90.09%
67%|          | 4/6 [05:16<02:37, 78.95s/it]

>> Epoch 4, Validation loss: 0.8321, Validation accuracy: 76.10%
[Epoch 5, Iteration 100] Training loss: 0.1828, Training accuracy: 94.53%
[Epoch 5, Iteration 200] Training loss: 0.2082, Training accuracy: 93.66%
[Epoch 5, Iteration 300] Training loss: 0.1924, Training accuracy: 94.59%
83%|          | 5/6 [06:34<01:18, 78.82s/it]

>> Epoch 5, Validation loss: 0.8452, Validation accuracy: 76.77%
[Epoch 6, Iteration 100] Training loss: 0.1146, Training accuracy: 96.78%
[Epoch 6, Iteration 200] Training loss: 0.0992, Training accuracy: 97.78%
[Epoch 6, Iteration 300] Training loss: 0.1146, Training accuracy: 96.44%
100%|         | 6/6 [07:53<00:00, 78.92s/it]

>> Epoch 6, Validation loss: 0.8404, Validation accuracy: 77.97%

Running experiment with lr=0.008, batch_size=64, epochs=10
Using SGD with momentum: 0.4
Unfreezing 10 layers
DataLoaders created successfully
Freeze all but the last n layers

0%|          | 0/10 [00:00<?, ?it/s]

```

```

[Epoch 1, Iteration 100] Training loss: 2.6135, Training accuracy: 22.03%
10%|          | 1/10 [01:15<11:20, 75.60s/it]

>> Epoch 1, Validation loss: 1.5058, Validation accuracy: 59.00%
[Epoch 2, Iteration 100] Training loss: 1.5593, Training accuracy: 52.73%
20%|          | 2/10 [02:31<10:04, 75.53s/it]

>> Epoch 2, Validation loss: 1.2078, Validation accuracy: 64.37%
[Epoch 3, Iteration 100] Training loss: 1.2642, Training accuracy: 62.41%
30%|          | 3/10 [03:46<08:48, 75.51s/it]

>> Epoch 3, Validation loss: 1.1116, Validation accuracy: 65.67%
[Epoch 4, Iteration 100] Training loss: 1.1450, Training accuracy: 65.64%
40%|          | 4/10 [05:02<07:34, 75.71s/it]

>> Epoch 4, Validation loss: 1.0282, Validation accuracy: 68.13%
[Epoch 5, Iteration 100] Training loss: 1.0090, Training accuracy: 69.62%
50%|          | 5/10 [06:18<06:18, 75.74s/it]

>> Epoch 5, Validation loss: 0.9850, Validation accuracy: 69.33%
[Epoch 6, Iteration 100] Training loss: 0.9583, Training accuracy: 70.61%
60%|          | 6/10 [07:33<05:02, 75.54s/it]

>> Epoch 6, Validation loss: 0.9519, Validation accuracy: 70.70%
[Epoch 7, Iteration 100] Training loss: 0.8780, Training accuracy: 73.39%
70%|          | 7/10 [08:49<03:46, 75.56s/it]

>> Epoch 7, Validation loss: 0.9337, Validation accuracy: 71.17%
[Epoch 8, Iteration 100] Training loss: 0.8095, Training accuracy: 74.19%
80%|          | 8/10 [10:04<02:31, 75.65s/it]

>> Epoch 8, Validation loss: 0.9183, Validation accuracy: 71.73%
[Epoch 9, Iteration 100] Training loss: 0.7882, Training accuracy: 75.97%
90%|          | 9/10 [11:20<01:15, 75.67s/it]

>> Epoch 9, Validation loss: 0.8955, Validation accuracy: 72.13%
[Epoch 10, Iteration 100] Training loss: 0.6991, Training accuracy: 78.16%
100%|         | 10/10 [12:36<00:00, 75.65s/it]

```



>> Epoch 10, Validation loss: 0.9016, Validation accuracy: 71.83%

Running experiment with lr=0.001, batch\_size=64, epochs=10

Using Adam

Unfreezing 10 layers

DataLoaders created successfully

Freeze all but the last n layers

0%| | 0/10 [00:00<?, ?it/s]

[Epoch 1, Iteration 100] Training loss: 2.7635, Training accuracy: 17.72%

10%| | 1/10 [01:15<11:21, 75.72s/it]

>> Epoch 1, Validation loss: 1.7767, Validation accuracy: 53.27%

[Epoch 2, Iteration 100] Training loss: 1.9063, Training accuracy: 43.20%

20%| | 2/10 [02:32<10:10, 76.35s/it]

>> Epoch 2, Validation loss: 1.5032, Validation accuracy: 58.63%

[Epoch 3, Iteration 100] Training loss: 1.7256, Training accuracy: 48.59%

30%| | 3/10 [03:48<08:53, 76.18s/it]

>> Epoch 3, Validation loss: 1.4154, Validation accuracy: 59.03%

[Epoch 4, Iteration 100] Training loss: 1.6510, Training accuracy: 50.23%

40%| | 4/10 [05:03<07:34, 75.82s/it]

>> Epoch 4, Validation loss: 1.3584, Validation accuracy: 60.33%

[Epoch 5, Iteration 100] Training loss: 1.6112, Training accuracy: 52.02%

50%| | 5/10 [06:19<06:18, 75.63s/it]

>> Epoch 5, Validation loss: 1.3105, Validation accuracy: 62.07%

[Epoch 6, Iteration 100] Training loss: 1.5860, Training accuracy: 52.23%

60%| | 6/10 [07:34<05:01, 75.49s/it]

>> Epoch 6, Validation loss: 1.3156, Validation accuracy: 61.23%

[Epoch 7, Iteration 100] Training loss: 1.6233, Training accuracy: 50.80%

70%| | 7/10 [08:50<03:46, 75.57s/it]

>> Epoch 7, Validation loss: 1.2991, Validation accuracy: 61.70%

[Epoch 8, Iteration 100] Training loss: 1.6438, Training accuracy: 50.94%

80%| | 8/10 [10:05<02:31, 75.64s/it]

```
>> Epoch 8, Validation loss: 1.3021, Validation accuracy: 62.10%  
[Epoch 9, Iteration 100] Training loss: 1.6038, Training accuracy: 52.34%
```


90%| | 9/10 [11:22<01:15, 75.88s/it]

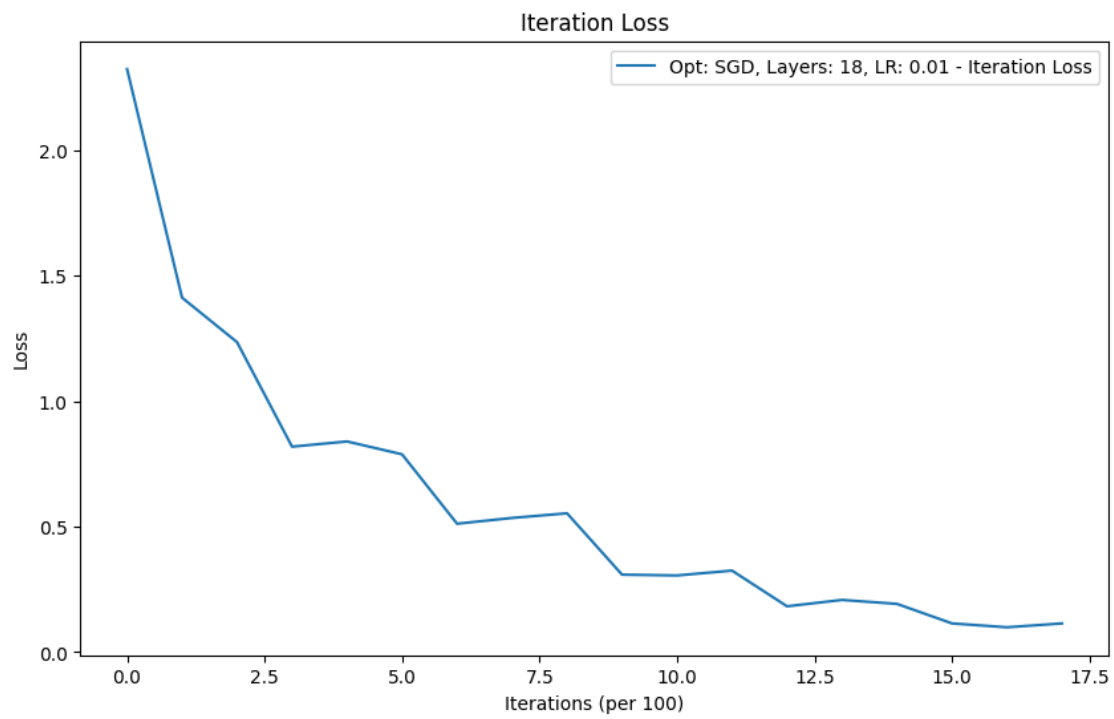

```
>> Epoch 9, Validation loss: 1.2857, Validation accuracy: 62.50%  
[Epoch 10, Iteration 100] Training loss: 1.6115, Training accuracy: 51.03%
```

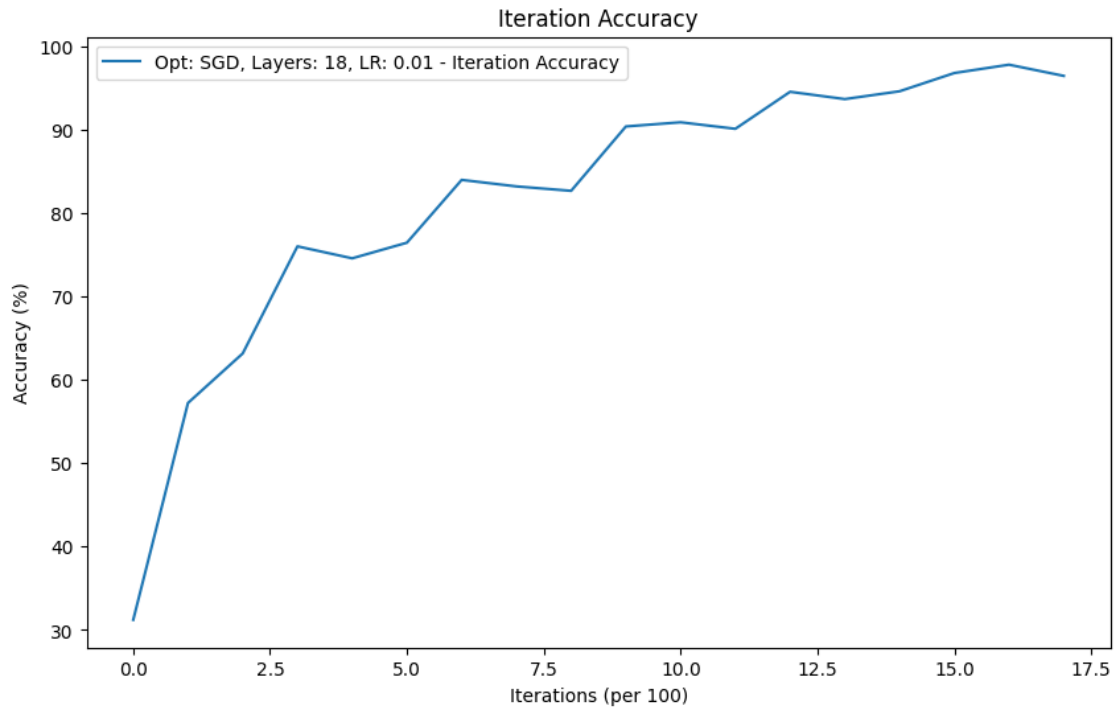
100%| | 10/10 [12:39<00:00, 75.92s/it]

```
>> Epoch 10, Validation loss: 1.2790, Validation accuracy: 61.70%
```


```
[ ]: plot_combined_experiment_results([experiment_results_3[0]])
```

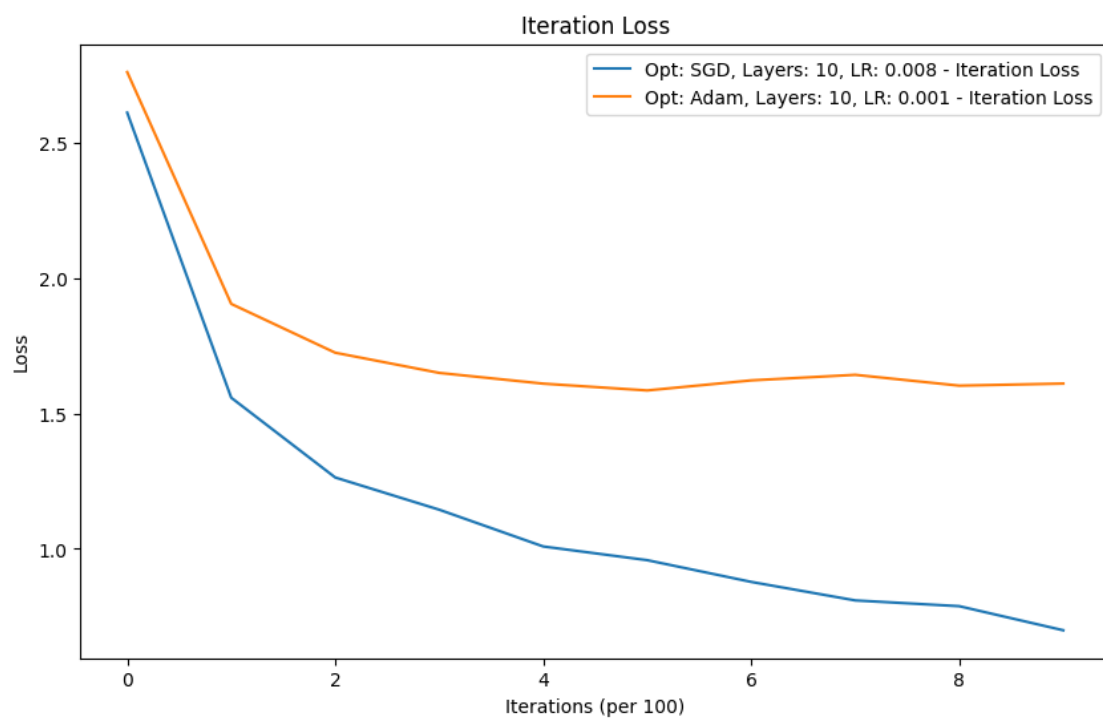



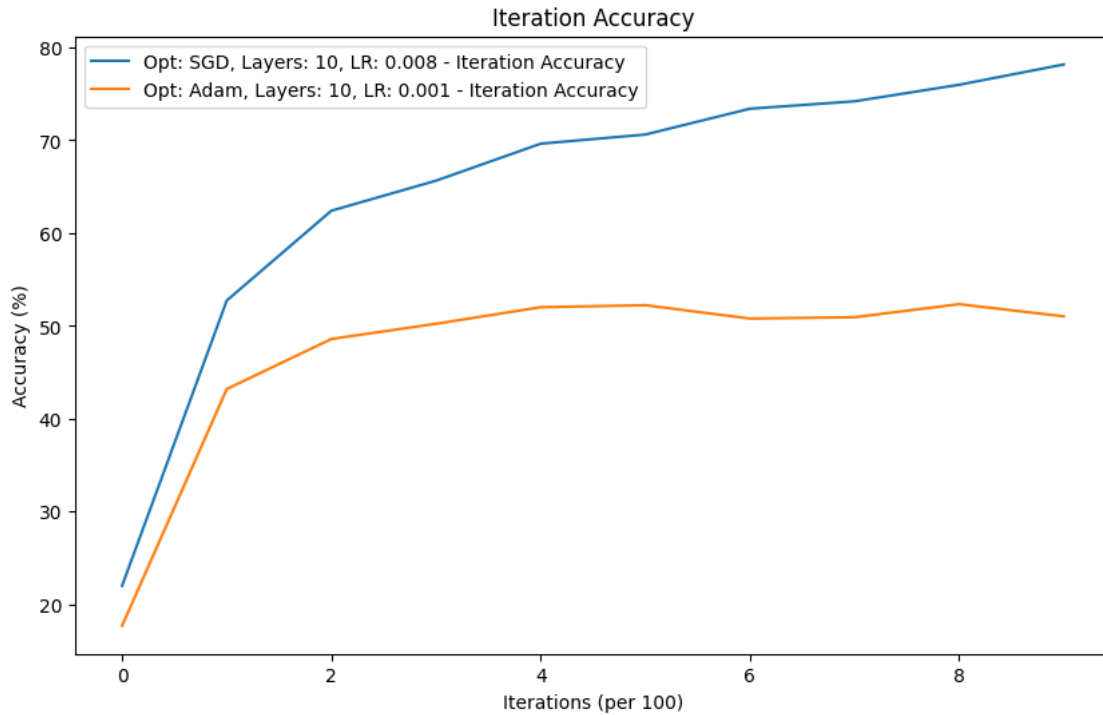




```
[ ]: model_results = [experiment_results_3[1]] + [experiment_results_3[2]]
plot_combined_experiment_results(model_results)
```







#### 4.4.1 Verdict/Insights

- Unable to beat the previous model.
- Increasing epochs for the first model is resulting in high training accuracy/low training loss, but stagnant validation accuracy/loss. This shows the model is memorizing the training data.
- SGD is a real difference maker here, Adam is performing very poorly.

## 5 Evaluation of the Best Model

- Using SGD with momentum: 0.5
- Unfreezing 18 layers
- 3 epochs
- Batch size of 32
- Learning rate of 0.01

### 5.1 Metrics for the Best Model (Validation Set)

```
[ ]: from sklearn.metrics import accuracy_score, precision_recall_fscore_support

def calculate_metrics(model, dataloader, device):
    model.eval()
    all_preds = []
    all_labels = []
```

```

with torch.no_grad():
    for data in dataloader:
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        _, preds = torch.max(outputs, 1)
        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

accuracy = accuracy_score(all_labels, all_preds)
precision, recall, f1, _ = precision_recall_fscore_support(all_labels,
↪all_preds, average='weighted')

print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")

return all_preds, all_labels

```

```

[ ]: model = initialize_model(18)
model.load_state_dict(torch.load('best_model.pth'))
model = model.to(device)

```

Unfreeze all the layers

```

[ ]: all_preds, all_labels = calculate_metrics(model, valloader, device)

```

```

Accuracy: 0.7687
Precision: 0.7757
Recall: 0.7687
F1-Score: 0.7669

```

## 5.2 Visualizing Predictions and Misclassifications (Validation Set)

```

[ ]: import matplotlib.pyplot as plt
import numpy as np

def denormalize(tensor, mean, std):
    tensor = tensor.clone() # Avoid modifying the original tensor
    for t, m, s in zip(tensor, mean, std):
        t.mul_(s).add_(m)
    return tensor

def visualize_predictions(model, dataloader, device, classes, num_images=10):
    model.eval()

```

```

images_so_far = 0
fig = plt.figure(figsize=(15, 15))

with torch.no_grad():
    for data in dataloader:
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        _, preds = torch.max(outputs, 1)

        for j in range(inputs.size()[0]):
            images_so_far += 1
            ax = plt.subplot(num_images // 2, 2, images_so_far)
            ax.axis('off')
            ax.set_title(f'Predicted: {classes[preds[j]]}, Actual: ↵
↵{classes[labels[j]]}')
            img = inputs.cpu().data[j].numpy().transpose((1, 2, 0))
            mean = np.array([0.485, 0.456, 0.406])
            std = np.array([0.229, 0.224, 0.225])
            img = std * img + mean
            img = np.clip(img, 0, 1)
            plt.imshow(img)

            if images_so_far == num_images:
                return

```

```

[ ]: classes = classes_to_keep
visualize_predictions(model, valloader, device, classes, num_images=10)

```



Predicted: apple\_pie, Actual: apple\_pie



Predicted: caesar\_salad, Actual: caesar\_salad



Predicted: bread\_pudding, Actual: carrot\_cake




Predicted: breakfast\_burrito, Actual: breakfast\_burrito



Predicted: bibimbap, Actual: bread\_pudding



Predicted: ceviche, Actual: ceviche



Predicted: apple\_pie, Actual: apple\_pie




Predicted: cannoli, Actual: cannoli



Predicted: caprese\_salad, Actual: caprese\_salad



Predicted: beet\_salad, Actual: caprese\_salad



## 6 Final Model Testing (Test Data)

```
[ ]: all_preds_test, all_labels_test = calculate_metrics(model, testloader, device)
```

Accuracy: 0.8005  
Precision: 0.8056  
Recall: 0.8005  
F1-Score: 0.7979

```
[ ]: visualize_predictions(model, testloader, device, classes, num_images=10)
```

Predicted: beef\_carpaccio, Actual: beef\_carpaccio



Predicted: carrot\_cake, Actual: breakfast\_burrito



Predicted: cheese\_plate, Actual: cheese\_plate



Predicted: beignets, Actual: beignets



Predicted: bibimbap, Actual: bibimbap



Predicted: beef\_tartare, Actual: beef\_tartare



Predicted: bruschetta, Actual: caprese\_salad



Predicted: breakfast\_burrito, Actual: breakfast\_burrito



Predicted: ceviche, Actual: ceviche



Predicted: bread\_pudding, Actual: bread\_pudding

